

HW3_report

November 26, 2020

Name: Chaitanya Sharadchandra Patil
PID: A53311586

Academic Integrity Policy: Integrity of scholarship is essential for an academic community. The University expects that both faculty and students will honor this principle and in so doing protect the validity of University intellectual work. For students, this means that all academic work will be done by the individual to whom it is assigned, without unauthorized aid of any kind. By including this in my report, I agree to abide by the Academic Integrity Policy mentioned above.

```
[139]: from tqdm import tqdm
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.image as image
import cv2
import skimage.morphology as skim_morph
import scipy.ndimage.measurements as scipy_im_measure
import pandas as pd
from scipy import signal
import math
```

0.0.1 Question 1

```
[140]: def canny_detector(im, thresh):
    im = im/255

    # smoothing
    smooth_filt_list = [[2,4,5,4,2],
                        [4,9,12,9,4],
                        [5,12,15,12,5],
                        [4,9,12,9,4],
                        [2,4,5,4,2]]

    smooth_filt = 1/159 * np.asarray(smooth_filt_list)
    smoothed = cv2.filter2D(im,-1,smooth_filt) # what exactly -1 does idk. Ask_
    →TA.

    # find gradients
```

```

k_x_list = [[-1,0,1],
            [-2,0,2],
            [-1,0,1]]

k_y_list = [[-1,-2,-1],
            [0,0,0],
            [1,2,1]]

k_x = np.asarray(k_x_list)
k_y = np.asarray(k_y_list)

# we need to convolve image with k_x and k_y
G_x = cv2.filter2D(im,-1,k_x)
G_y = cv2.filter2D(im,-1,k_y)

# gradient magnitude
gradient_image = np.sqrt(G_x**2 + G_y**2)
plt.figure(figsize = [10,10])
plt.imshow(gradient_image)
plt.title('Original Gradient Magnitude Image')
# gradient angle
G_x = np.where(G_x == 0, 0.00000001, G_x)
cont_angles = np.arctan(G_y/G_x)

# NMS

# rounding angles
disc_angles = np.zeros(cont_angles.shape)
disc_angles = np.where(cont_angles<(-3*math.pi/8),(-math.pi/2),cont_angles)
disc_angles = np.where((cont_angles>=(-3*math.pi)/8) & (cont_angles<(-math.
→pi/8)), -math.pi/4, disc_angles)
disc_angles = np.where((cont_angles>=-math.pi/8) & (cont_angles<math.pi/
→8), 0, disc_angles)
disc_angles = np.where((cont_angles>=math.pi/8) & (cont_angles<3*math.pi/
→8), math.pi/4, disc_angles)
disc_angles = np.where(cont_angles>=3*math.pi/8 ,math.pi/2, disc_angles)

# pad mirroring
gradient_padded = np.pad(gradient_image, 1, mode = 'symmetric')
disc_angle_padded = np.pad(disc_angles, 1, mode = 'symmetric')

supressed = np.zeros(gradient_padded.shape)

for i in range(1,supressed.shape[0]-1):
    row_num = i

```

```

for j in range(1,supressed.shape[1]-1):
    col_num = j

    # above is correct 100%

    angle = disc_angle_padded[row_num, col_num]
    magnitude = gradient_padded[row_num, col_num]

    if(angle == 0):
        if((magnitude<gradient_padded[row_num, col_num-1]) or
→(magnitude<gradient_padded[row_num, col_num+1])):
            supressed[row_num, col_num] = 0
        else:
            supressed[row_num, col_num] = gradient_padded[row_num,
→col_num]

    if(angle == math.pi/2 or angle == -math.pi/2):
        if(magnitude<gradient_padded[row_num-1, col_num] or
→magnitude<gradient_padded[row_num+1, col_num]):
            supressed[row_num, col_num] = 0
        else:
            supressed[row_num, col_num] = gradient_padded[row_num,
→col_num]

    if(angle == math.pi/4):
        if(magnitude<gradient_padded[row_num-1, col_num+1] or
→magnitude<gradient_padded[row_num+1, col_num-1]):
            supressed[row_num, col_num] = 0
        else:
            supressed[row_num, col_num] = gradient_padded[row_num,
→col_num]

    if(angle == -math.pi/4):
        if(magnitude<gradient_padded[row_num+1, col_num+1] or
→magnitude<gradient_padded[row_num-1, col_num-1]):
            supressed[row_num, col_num] = 0
        else:
            supressed[row_num, col_num] = gradient_padded[row_num,
→col_num]


plt.figure(figsize = [10,10])
plt.imshow(supressed)
plt.title('Image After NMS')
# thresholding
thresholded_supressed = np.where(supressed<thresh,0,1)


return(thresholded_supressed)

```

```
[141]: g_sel_color = cv2.imread('geisel.jpg')  
g_sel = cv2.cvtColor(g_sel_color, cv2.COLOR_BGR2GRAY)
```


```
[142]: canny_image_geisel = canny_detector(g_sel, 0.9)
```





```
[143]: plt.figure(figsize = [10,10])  
plt.imshow(canny_image_geisel)  
plt.title('Final edge image after thresholding')
```

```
[143]: Text(0.5, 1.0, 'Final edge image after thresholding')
```




The value of t_e used was 0.9.

0.0.2 Question 2

Part (i)

```
[184]: # reading image and displaying unpadded original image
im_car = image.imread('Car.tif')

fig = plt.figure(figsize = [7.5,7.5])
ax = fig.add_subplot()
mpbl = ax.imshow(im_car, cmap = 'gray')
ax.set_title('Unpadded Original Image');
ax.set_xticks([])
ax.set_yticks([])
fig.colorbar(mpbl);
```



```
[197]: # normalizing and padding
# im_car = im_car/255
im_car_padded = np.pad(im_car, ((133,), (172,)))
```

```
[198]: # taking fft
fft_im_car_padded = np.fft.fft2(im_car_padded)
```


```
[199]: # plotting 2D DFT log_magnitude image
fig = plt.figure(figsize = [7.5,7.5])
ax = fig.add_subplot()
mpbl = ax.imshow(np.log(abs(np.fft.fftshift(fft_im_car_padded))))
```



```

ax.set_xlabel('u')
ax.set_ylabel('v')
# ax.set_xticks([])
# ax.set_yticks([])
ax.set_title('2D DFT log magnitude image of car')
fig.colorbar(mpl);

```



After zooming into the image using the mpld3 package, I have found the indices of the maximae to the left of the image to be (row_number, column_number):

346, 336

346, 421

339, 168

339, 83

We will store this in a list for easy access.

```
[200]: notch_reject_locations = [[346,336], [346,421], [339,168], [339, 83]]
```

```
[201]: x_axis = np.linspace(-256,255,512)
y_axis = np.linspace(-256,255,512)
[u,v] = np.meshgrid(x_axis,y_axis)
```

We will write the equation for one notch reject filter.

```
[202]: def dk(u, v, uk, vk):
        return np.sqrt(((u - uk)**2 + (v - vk)**2))
```

```
[203]: def d_neg_k(u, v, uk, vk):
        return np.sqrt(((u + uk)**2 + (v + vk)**2))
```

```
[204]: def hnr_term(D_0, n, uk, vk):

        hnr = (1/(1+((D_0/dk(u,v,uk,vk))**(2*n))))*(1/(1+((D_0/
→d_neg_k(u,v,uk,vk))**(2*n))))

        return(hnr)
```

```
[205]: hnr_car = 1

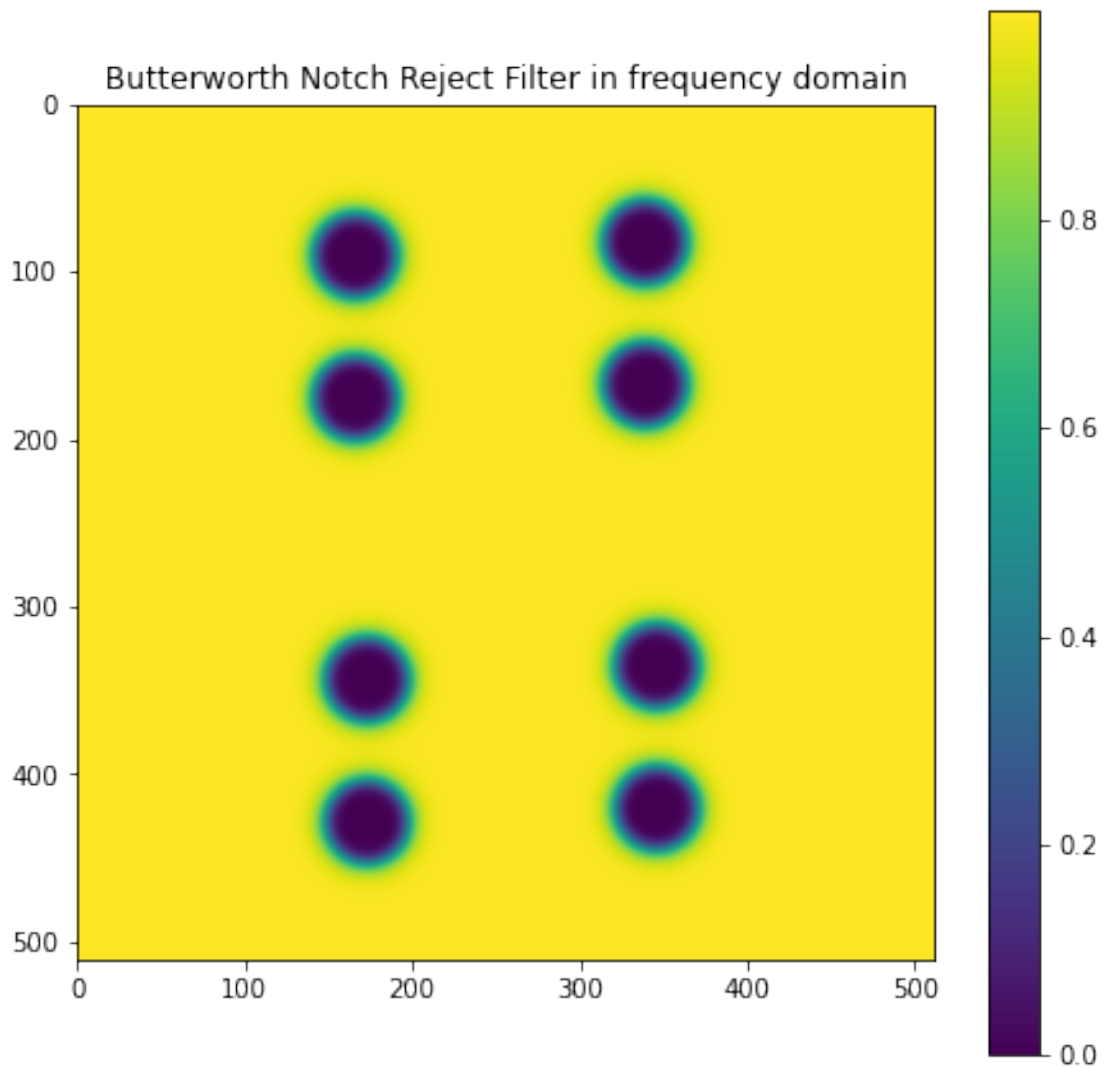
for i in range(len(notch_reject_locations)):

    uk = notch_reject_locations[i][0] - 256
    vk = notch_reject_locations[i][1] - 256
    n = 4
    D_0 = 25
    hnr_car = hnr_car * hnr_term(D_0, n, uk, vk)
```

/home/chaitanya/Documents/odas_test_notebooks/odas/lib/python3.6/site-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in true_divide

This is separate from the ipykernel package so we can avoid doing imports until

```
[206]: # plotting butterworth Notch Reject Filter in frequency domain
fig = plt.figure(figsize=[7.5,7.5])
ax = fig.add_subplot()
mpbl = ax.imshow(hnr_car)
ax.set_title('Butterworth Notch Reject Filter in frequency domain')
fig.colorbar(mpbl);
```



```
[207]: # filtering by multiplication with filter in f domain  
  
filtered_fft = hnr_car*np.fft.fftshift(fft_im_car_padded)  
  
# taking ifftshift  
  
fft_inv_shift = np.fft.ifftshift(filtered_fft)
```


```
[209]: # taking inverse 2D ffit  
  
inv_fft = np.fft.ifft2(fft_inv_shift)  
  
# taking abs of the resulting complex vector
```

```
filtered_padded_image = abs(inv_fft)

# cropping image appropriately
filtered_image = filtered_padded_image[133:378, 172:339]
```

```
[210]: # plotting final filtered image

fig = plt.figure(figsize=[7.5,7.5])
ax = fig.add_subplot()
mpbl = ax.imshow(abs(filtered_image), cmap = 'gray')
ax.set_title('Final Filtered Image')
fig.colorbar(mpbl);
```



Parameters used were:

`n = 4`

`D_0 = 25 u1 = 346`

`v1 = 336`

`u2 = 346`

`v2 = 421`

`u3 = 339`

`v3 = 168`

`u4 = 339`

`v4 = 83`

Part (ii)

```
[211]: street = image.imread('Street.png')
```

```
[212]: # reading image and displaying unpadded original image
fig = plt.figure(figsize = [7.5,7.5])
ax = fig.add_subplot()
mpbl = ax.imshow(street, cmap = 'gray')
ax.set_title('Unpadded Original Image');
ax.set_xticks([])
ax.set_yticks([])
fig.colorbar(mpbl);
```



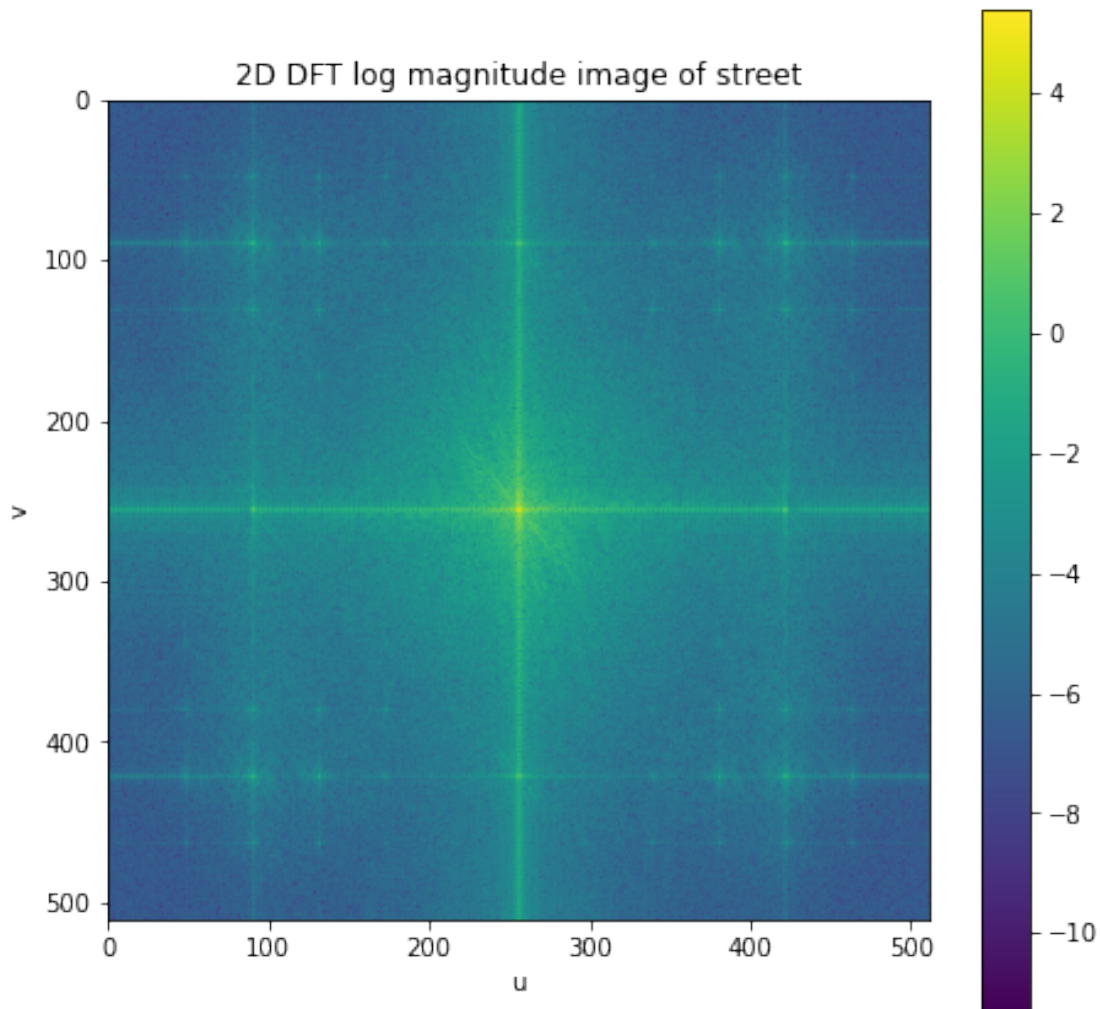
```
[213]: street = street[:,1:]
```

```
[214]: im_street_padded = np.pad(street, (((512-street.shape[0])/2,)), ((512-street.  
→shape[1])/2,)))  
im_street_padded = im_street_padded/255
```

```
[215]: fft_im_street_padded = np.fft.fft2(im_street_padded)
```

```
[216]: # plotting 2D DFT log_magnitude image  
fig = plt.figure(figsize = [7.5,7.5])  
ax = fig.add_subplot()  
mpbl = ax.imshow(np.log(abs(np.fft.fftshift(fft_im_street_padded))))  
ax.set_xlabel('u')
```

```
ax.set_ylabel('v')
# ax.set_xticks([])
# ax.set_yticks([])
ax.set_title('2D DFT log magnitude image of street')
fig.colorbar(mpl);
```



```
[220]: %matplotlib inline
import mpld3
mpld3.enable_notebook()
```

We want to remove the burst along $u = 0$ axis and $v = 0$ axis

We will zoom into the image to extract the indices of the bursts:

burst along $u = 0$ axis.

422, 256

burst along $v = 0$ axis.

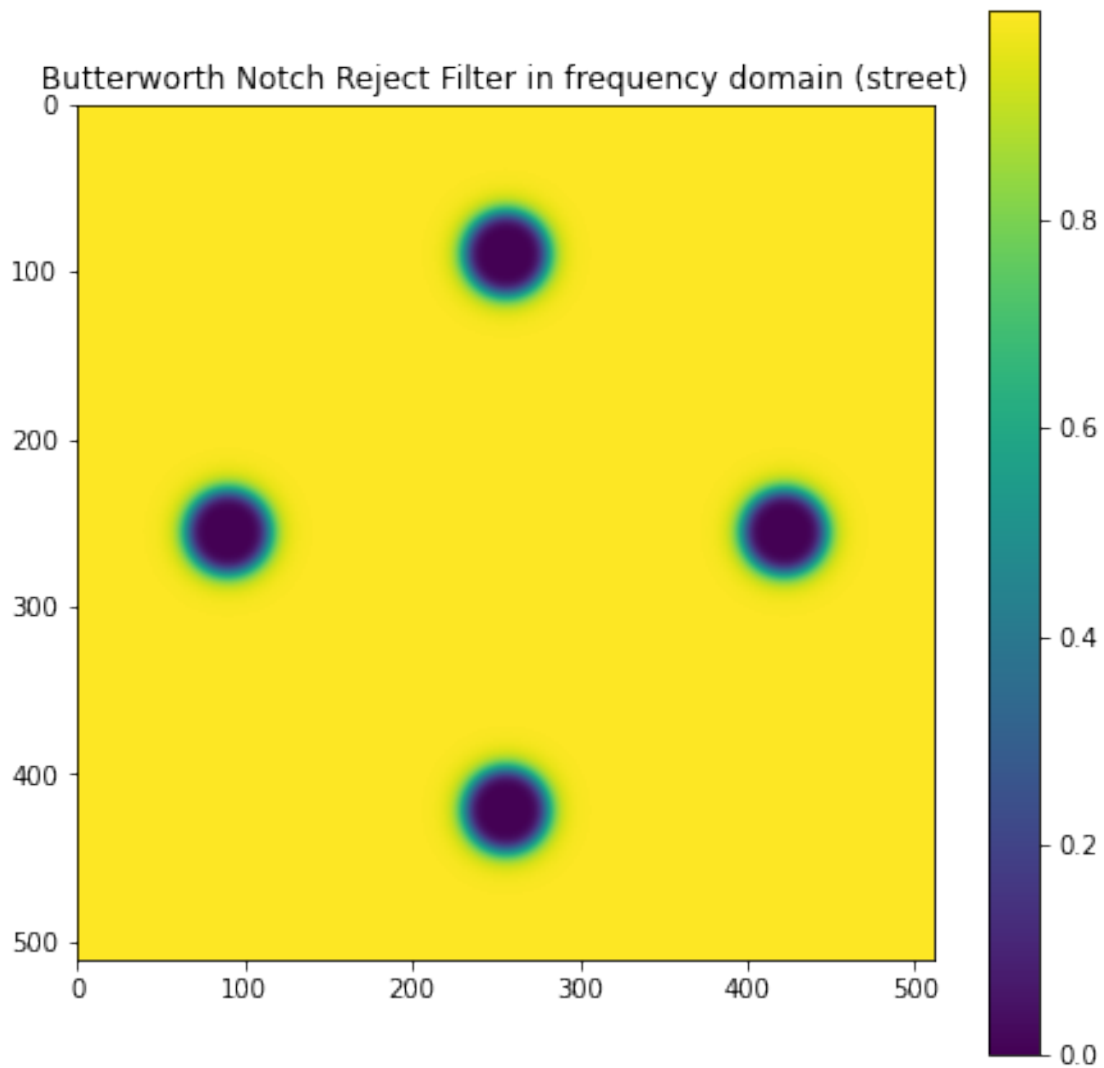
256, 422

```
[164]: term_21 = hnr_term(25, 4, 0, 422-256)
term_22 = hnr_term(25, 4, 422-256, 0)
term_2 = term_21*term_22
```

/home/chaitanya/Documents/odas_test_notebooks/odas/lib/python3.6/site-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in true_divide

This is separate from the ipykernel package so we can avoid doing imports until

```
[165]: # plotting butterworth Notch Reject Filter in frequency domain
fig = plt.figure(figsize=[7.5,7.5])
ax = fig.add_subplot()
mpbl = ax.imshow(term_2)
ax.set_title('Butterworth Notch Reject Filter in frequency domain (street)')
fig.colorbar(mpbl);
```

```
[166]: # multiplying fft with this mask

filtered_fft = np.fft.fftshift(fft_im_street_padded) * term_2

# we need to shift this filtered_fft in the inverse direction of the shift
filtered_shifted_fft = np.fft.fftshift(filtered_fft)

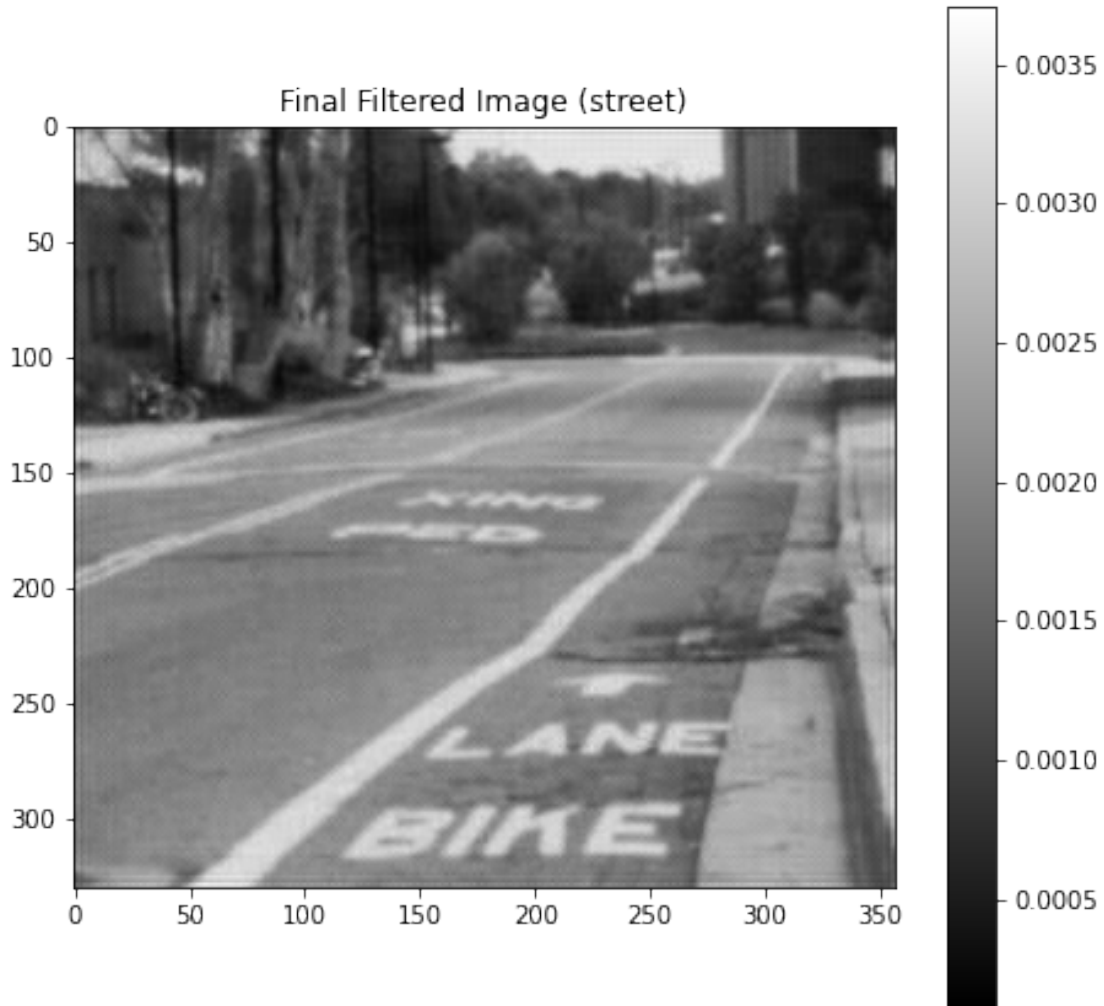
# we need to take the inverse fft of this to get the padded filtered image
padded_filtered_image = np.fft.ifft2(filtered_shifted_fft)

# we usually get complex array due to some inaccuracies. We will take only the
→ magnitude of the array.
```

```
filtered_image = padded_filtered_image[90:421,77:434]
```

```
[167]: # plotting final filtered image

fig = plt.figure(figsize=[7.5,7.5])
ax = fig.add_subplot()
mpbl = ax.imshow(abs(filtered_image), cmap = 'gray')
ax.set_title('Final Filtered Image (street)')
fig.colorbar(mpbl);
```



Parameters used were: $n = 4$

$D_0 = 25$

$u_1 = 422$

$v_1 = 256$

$u_2 = 256$

$v_2 = 422$

0.0.3 Question 3

Acknowledgement: Much of the code below is from https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

All of the code in the part until after training is 100% from https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

Specific references have been given when applicable.

```
[1]: import torch
import torchvision
import torchvision.transforms as transforms
```

```
[2]: # Code from https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                          shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified

Files already downloaded and verified

```
[3]: # Code from https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

import matplotlib.pyplot as plt
import numpy as np

# functions to show an image

def imshow(img):
    img = img / 2 + 0.5     # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
```

```
plt.show()

# get some random training images
dataiter = iter(trainloader)
images, labels = dataiter.next()

# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print(' '.join('%5s' % classes[labels[j]] for j in range(4)))
```



horse car bird deer

[4]: # Code from https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
```

```

x = F.relu(self.fc1(x))
x = F.relu(self.fc2(x))
x = self.fc3(x)
return x

```

```
net = Net()
```

[5]: *# Code from https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html*

```
import torch.optim as optim
```

```
criterion = nn.CrossEntropyLoss()
```

```
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

[6]: *# Code from https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html*

```
for epoch in range(2): # loop over the dataset multiple times
```

```
    running_loss = 0.0
```

```
    for i, data in enumerate(trainloader, 0):
```

```
        # get the inputs; data is a list of [inputs, labels]
```

```
        inputs, labels = data
```

```
        # zero the parameter gradients
```

```
        optimizer.zero_grad()
```

```
        # forward + backward + optimize
```

```
        outputs = net(inputs)
```

```
        loss = criterion(outputs, labels)
```

```
        loss.backward()
```

```
        optimizer.step()
```

```
        # print statistics
```

```
        running_loss += loss.item()
```

```
    if i % 2000 == 1999: # print every 2000 mini-batches
```

```
        print('%d, %5d loss: %.3f' %
```

```
              (epoch + 1, i + 1, running_loss / 2000))
```

```
        running_loss = 0.0
```

```
print('Finished Training')
```

/home/chaitanya/Documents/odas_test_notebooks/odas/lib/python3.6/site-packages/torch/autograd/__init__.py:132: UserWarning: CUDA initialization: Found no NVIDIA driver on your system. Please check that you have an NVIDIA GPU and

installed a driver from <http://www.nvidia.com/Download/index.aspx> (Triggered internally at /pytorch/c10/cuda/CUDAFuncions.cpp:100.)
allow_unreachable=True) # allow_unreachable flag

```
[1, 2000] loss: 2.229
[1, 4000] loss: 1.866
[1, 6000] loss: 1.675
[1, 8000] loss: 1.556
[1, 10000] loss: 1.485
[1, 12000] loss: 1.446
[2, 2000] loss: 1.364
[2, 4000] loss: 1.351
[2, 6000] loss: 1.343
[2, 8000] loss: 1.301
[2, 10000] loss: 1.288
[2, 12000] loss: 1.269
Finished Training
```

ii. How many images and batches are used to train the network?

There are 50000 images.

Each batch has 4 images, so the total number of batches is 12500, since $12500 * 4 = 50000$.

```
[12]: trainset.__len__()
```

```
[12]: 50000
```

iii. Do we normalize the images? What do we do in the example?

We do normalize the images. In the example, we defined a transform as consisting of toTensor and normalization. In the normalization, we normalize the images.

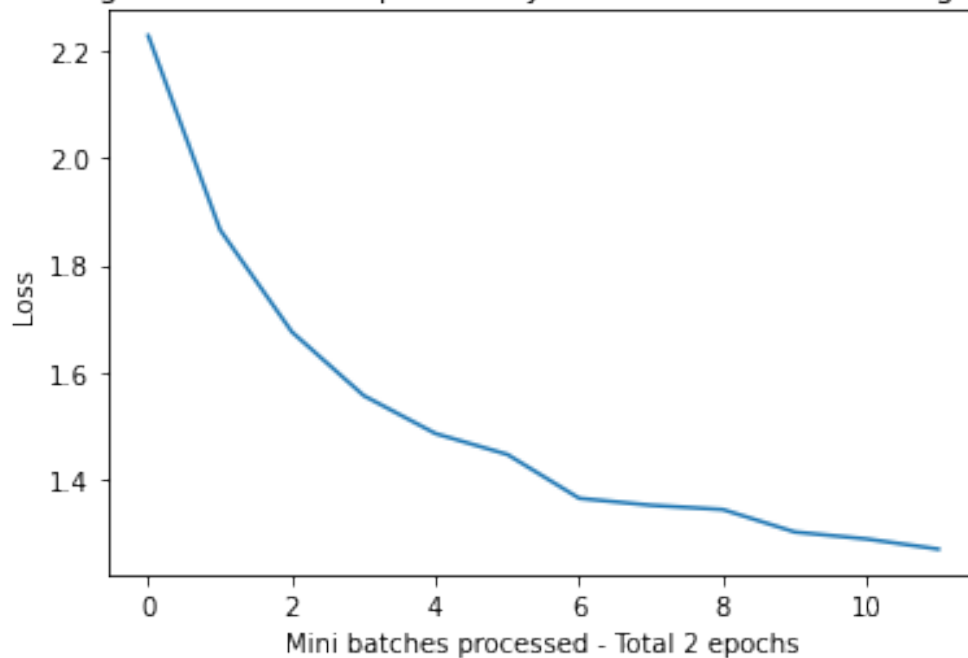
iv. Plotting training loss

```
[7]: # plotting training loss
```

```
losses = [2.229, 1.866, 1.675, 1.556, 1.485, 1.446, 1.364, 1.351, 1.343, 1.301,
→1.288, 1.269]
```

```
[10]: fig = plt.figure()
ax = fig.add_subplot()
ax.plot(losses)
ax.set_xlabel('Mini batches processed - Total 2 epochs')
ax.set_ylabel('Loss')
ax.set_title('Plotting loss values sampled every 2000 mini batches during
→training');
```

Plotting loss values sampled every 2000 mini batches during training



- v. Now the network is done training. Can you check some successful cases and some failure cases (show some images classified by the network)?

```
[52]: # Code adopted from https://pytorch.org/tutorials/beginner/blitz/
      ↪cifar10_tutorial.html

      # to check the results of the network on the images, we must first retrieve
      ↪images
      # we should do so from the test loader

      loader_iterable = iter(testloader)
      images, labels = loader_iterable.next() # this should retrieve a minibatch

      # we must propagate our images through the network.

      outputs = net(images)
```

```
[53]: # Code adopted from https://pytorch.org/tutorials/beginner/blitz/
      ↪cifar10_tutorial.html

      imshow(torchvision.utils.make_grid(images))
      print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))

      _, predicted = torch.max(outputs, 1)
```



```
print('Predicted: ', ' '.join('%5s' % classes[predicted[j]] for j in range(4)))
```



```
GroundTruth:  cat  ship  ship plane
Predicted:    cat  ship  ship  ship
```

Above, we get to see that for the first three images, we got the right output, but this is not the case for the last image.

vi. Can you visualize the output of the 1st layer of CNN using one image from the training set?

```
[97]: # the kernel has been trained. All we need to do is convolve with a few images.

# step 1: obtain images

# we will use our images from previous step
```

```
[98]: # our weights are stored in net
# https://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html,
→ from where I have adopted this code

all_net_params = list(net.parameters())
```

what if we defined a sub_net using just the first layer params of net?

If we want to do that, we will need to create a new class of network, I think

```
[99]: # Code adopted from https://pytorch.org/tutorials/beginner/blitz/
→ cifar10_tutorial.html

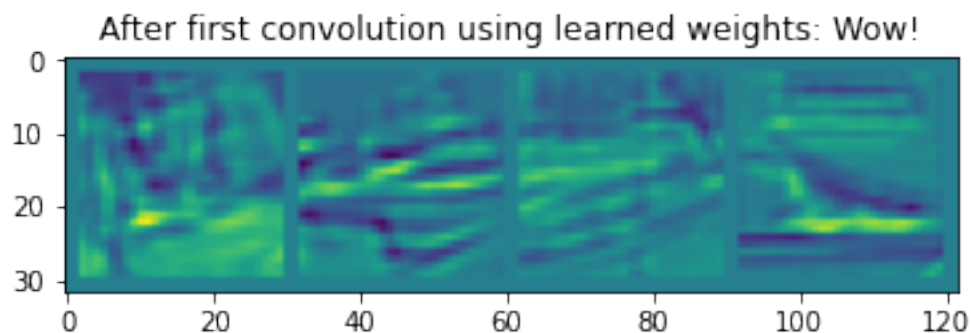
class Net2(nn.Module):
    def __init__(self, params_layer_1):
        super(Net2, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.conv1.weight = params_layer_1 # adopted from https://discuss.
→ pytorch.org/t/how-to-set-nn-conv2d-weights/67407
    def forward(self, x):
        x = self.conv1(x)
```

```
return x
```

```
[100]: net_2 = Net2(all_net_params[0])  
first_layer_output = net_2(images)
```

```
[118]: def imshow_detach(img):  
        """Adopted from the function imshow() at https://pytorch.org/tutorials/  
        ↳beginner/blitz/cifar10_tutorial.html"""  
        npimg = img.detach().numpy()  
        npimg_tp = np.transpose(npimg, (1, 2, 0))  
        npimg_tp_avg = np.mean(npimg_tp, axis = 2)  
        plt.imshow(npimg_tp_avg)  
        plt.title('After first convolution using learned weights: Wow!')  
        plt.show()
```

```
[119]: # Code adopted from https://pytorch.org/tutorials/beginner/blitz/  
        ↳cifar10_tutorial.html  
imshow(torchvision.utils.make_grid(images))  
imshow_detach(torchvision.utils.make_grid(first_layer_output))
```



Note: I averaged feature maps across 6 channels of the resultant feature map.

References:

1. <https://www.youtube.com/watch?v=sRFM5IEqR2w>
2. <https://matplotlib.org/3.1.0/gallery/mplot3d/surface3d.html>